

# FULLSTACK RUST

*The Complete Guide to Building Apps with the Rust  
Programming Language and Friends*

# Fullstack Rust

*The Complete Guide to Building Apps with the Rust Programming Language and Friends*

Written by Andrew Weiss

Edited by Nate Murray

© 2020 Fullstack.io

All rights reserved. No portion of the book manuscript may be reproduced, stored in a retrieval system, or transmitted in any form or by any means beyond the number of purchased copies, except for a single backup or archival copy. The code may be used freely in your projects, commercial or otherwise.

The authors and publisher have taken care in preparation of this book, but make no expressed or implied warranty of any kind and assume no responsibility for errors or omissions. No liability is assumed for incidental or consequential damages in connection with or arising out of the use of the information or programs contained herein.

Published by Fullstack.io.



# Contents

Book Revision . . . . .	1
Join Our Discord . . . . .	1
Bug Reports . . . . .	1
Be notified of updates via Twitter . . . . .	1
We'd love to hear from you! . . . . .	1
<b>Introduction . . . . .</b>	<b>1</b>
Why Rust? . . . . .	1
Why not Rust . . . . .	8
This book's mission . . . . .	9
Setting expectations based on your background . . . . .	10
Getting your environment setup . . . . .	12
Rustup . . . . .	13
Cargo . . . . .	13
IDEs, RLS, Editors . . . . .	14
Clippy . . . . .	14
Rustfmt . . . . .	15
Documentation . . . . .	15
The Nomicon . . . . .	16
<b>Summary . . . . .</b>	<b>17</b>
<b>Making Your First Rust App . . . . .</b>	<b>18</b>
Getting started . . . . .	18
Binary vs. library . . . . .	18
The generated project . . . . .	19
Crates . . . . .	21
Making our crate a library . . . . .	22

## CONTENTS

Trade-offs . . . . .	24
Print a list of numbers . . . . .	24
Testing our code . . . . .	39
Wrapping up . . . . .	40
<b>Making A Web App With Actix . . . . .</b>	<b>41</b>
Web Ecosystem . . . . .	41
Starting out . . . . .	44
Handling our first request . . . . .	51
<b>Adding State to Our Web App . . . . .</b>	<b>65</b>
Recap and overview . . . . .	65
Adding state . . . . .	65
Receiving input . . . . .	75
Custom error handling . . . . .	81
Handling path variables . . . . .	87
Wrapping up . . . . .	90
<b>Even More Web . . . . .</b>	<b>91</b>
Crates to know . . . . .	91
Building a blog . . . . .	92
Users . . . . .	96
Building the application . . . . .	100
Examples . . . . .	126
Extending our application . . . . .	129
Adding routes for posts . . . . .	136
Extending further: comments . . . . .	141
Adding routes for comments . . . . .	149
<b>Examples . . . . .</b>	<b>154</b>
Create a post . . . . .	154
Create a post . . . . .	154
Publish a post . . . . .	155
Comment on a post . . . . .	155
List all posts . . . . .	155
See posts . . . . .	157
Publish other post . . . . .	157

## CONTENTS

List all posts again . . . . .	158
See users comments . . . . .	159
See post comments . . . . .	160
Wrapping up . . . . .	161
<b>What is Web Assembly? . . . . .</b>	<b>162</b>
Intro to Web Assembly . . . . .	162
Rust in the browser . . . . .	164
The Smallest Wasm Library . . . . .	165
Working with primitives . . . . .	167
Working with complex types . . . . .	172
The Real Way to Write Wasm . . . . .	186
Other Wasm Topics . . . . .	191
<b>Command Line Applications . . . . .</b>	<b>193</b>
Initial setup . . . . .	193
Making an MVP . . . . .	194
Recap . . . . .	235
Adding a configuration file . . . . .	235
Adding sessions . . . . .	243
Syntax highlighting . . . . .	259
Summary . . . . .	266
<b>Macros . . . . .</b>	<b>267</b>
Overview . . . . .	267
Declarative Macros . . . . .	268
Procedural Macros . . . . .	274
Writing a custom derive . . . . .	276
Using our custom derive . . . . .	309
Wrapping up . . . . .	314
<b>Changelog . . . . .</b>	<b>315</b>
Revision 5 (02-20-2020) . . . . .	315
Revision 4 (02-19-2020) . . . . .	315
Revision 3 (01-29-2020) . . . . .	315
Revision 2 (11-25-2019) . . . . .	315
Revision 1 (10-29-2019) . . . . .	315

# Book Revision

Revision 5 - 2020-02-20

## Join Our Discord

Come chat with other readers of the book in the official newline Discord channel:

Join here: <https://newline.co/discord/rust><sup>1</sup>

## Bug Reports

If you'd like to report any bugs, typos, or suggestions just email us at: [us@fullstack.io](mailto:us@fullstack.io).

## Be notified of updates via Twitter

If you'd like to be notified of updates to the book on Twitter, follow us at [@fullstackio](https://twitter.com/fullstackio)<sup>2</sup>.

## We'd love to hear from you!

Did you like the book? Did you find it helpful? We'd love to add your face to our list of testimonials on the website! Email us at: [us@fullstack.io](mailto:us@fullstack.io)<sup>3</sup>.

---

<sup>1</sup><https://newline.co/discord/rust>

<sup>2</sup><https://twitter.com/fullstackio>

<sup>3</sup><mailto:us@fullstack.io>

# Introduction

There are numerous reasons to be hopeful about the future of computing, one of which is the existence and continued progression of the Rust programming language.

We are currently in the fifth era of programming language evolution. This is an era where languages have been able to take all of the learnings since the 1950s and incorporate the best parts into languages each with its own cohesive vision.

We have specialized languages cropping up for a wide variety of tasks and countless general purpose languages being actively developed and used. There are significant resources in industry to invest in language design and development which complement the vibrant academic community. With tools like LLVM and the explosion of open source, creating a language has never been easier.

It is in this environment that Rust has been voted the “most loved programming language” in the Stack Overflow Developer Survey every year since 2016. Standing out in this increasingly crowded world of languages is enough of a reason to ask why Rust?

## Why Rust?

There are a few potential readings of this question: why should I learn Rust, why are others using Rust, why should I choose Rust over language X? These are all relevant, but I want to start with a bit of a philosophical argument for Rust independent of these specific points.

There is a limit to how transformative an experience you can have when learning a language in a similar paradigm to one you already know. Every language and paradigm has an intrinsic style that is forced on you as you try to solve problems.

If you work within that style then your code will flow naturally and the language will feel like it is working with you. On the other hand, if you fight the natural style of the language you will find it hard or impossible to express your ideas.

Moreover, learning and working with a language will teach you ways to be more effective based on how the language guides you based on its natural design. How much you are able to learn is a function of how much your prior experience and mental models cover the new language.

Rust borrows a lot of ideas from other languages and is truly multi-paradigm, meaning you can write mostly functional code or mostly imperative code and still fit nicely within the language. The most unique feature of the language, the borrow checker, is a system that enforces certain invariants which allow you to make certain safety guarantees. Even this is built on prior art found in earlier languages.

All of these good ideas from the world of programming language design combine in a unique way to make Rust a language that truly makes you think about writing code from a novel perspective. It does not matter how much experience you have, learning Rust will forever change the way you write code for the better.

Okay with that philosophical argument out of the way, let's dig in to some specifics of why Rust is a exciting.

To help guide this discussion, we can break things down into a few broad categories.

## **On language comparisons**

There is no best programming language. Almost every task has a variety of languages which could be the right tool. Every language comes with good parts and bad parts. Evaluating these trade-offs when faced with a particular problem space is an art unto itself. Therefore, nothing in this book is intended to disparage or denigrate any particular alternative language. The primary goal of this book is to faithfully present Rust. That being said, sometimes comparisons with other languages are instructive and are meant to be instructive rather than as fuel in a flame war.

## **Language features**

There are a lot of features of Rust which make it a great tool for a great number of tasks. Some highlights include:

- Performance



- Strong, static, expressive type system
- Great error messages
- Modern generics
- Memory safety
- Fearless concurrency
- Cross platform
- C interoperability

Let's briefly go through some of these which are probably the biggest reasons that Rust gets talked about.

## Performance

Rust is exceptionally fast, in the same ballpark as C and C++. For some programs, specifically due to the lack of pointer aliasing, the Rust compiler can sometimes have enough information to optimize code to be faster than what is possible in C without directly writing assembly. For the vast majority of use cases, you should consider Rust to be fast enough.

Often the most obvious way to write a program is also the fastest. Part of this comes from the commitment to zero-cost abstractions, which are summarized by Bjarne Stroustrup, the creator of C++, as:

What you don't use, you don't pay for. And further: What you do use, you couldn't hand code any better.

Most of the abstractions in Rust, for example iterators, are zero-cost by this definition. The most efficient way to traverse a vector of data is to use a for loop which uses an iterator trait. The generated assembly is usually as good as you could hope for had you written it by hand.

The other aspect of performance is memory consumption. Rust does not have a garbage collector so you can use exactly as much memory as is strictly necessary at any given time. Due to the design of the language, you start to think and see every memory allocation. Using less memory is often easier than the converse. The rest of the language is designed around making working without a garbage collector painless.

## Type system

The type system of Rust is influenced by the long lineage of functional programming languages such as ML and Haskell. It is static, nominal, strong, and for the most part inferred. Don't worry if that didn't mean anything to you, but if it did then great. You encode the ideas and constraints of your problem with types. You only have to specify types in a few places with the rest able to be inferred.

A type system is often called expressive if it is easy to encode your ideas. Rust has powerful abstraction facilities like sum and product types, tuples, generics, etc. which put the type system definitely in the expressive camp.

There are some concepts which are impossible to express in static type systems. However, most things you encounter in practice are expressible in Rust's type system.

The compiler then checks everything for you so that you get faster feedback about potential problems. As a result, entire classes of bugs are impossible because of this static typing.

## Memory safety

A language is memory safe if certain classes of bugs related to memory access are not possible. Likewise, a language can be called memory unsafe if certain bugs are possible. A non-exhaustive list of memory related bugs include: dereferencing null pointers, use-after free, dangling pointers, buffer overflows.

If you have never written code in a memory unsafe language then these might sound like gibberish to you, which is fine. The important point is this class of bugs is a consistent and large source of security vulnerabilities in systems implemented with memory unsafe languages. For example, about [20% of CVEs<sup>4</sup>](https://www.cvedetails.com/product/47/Linux-Linux-Kernel.html?vendor_id=33) ever filed against the Linux kernel are due to memory corruption or overflows. Linux is implemented primarily in C, a spectacularly memory unsafe language.

Memory safety bugs are bad for security and reliability. They lead to vulnerabilities and they lead to crashes. If you can rule these out at compile time then you are in a much better state of the world.

Rust is designed to be memory safe, and thus it does not permit null pointers, dangling pointers, or data races in safe code. There are many interacting features

---

<sup>4</sup>[https://www.cvedetails.com/product/47/Linux-Linux-Kernel.html?vendor\\_id=33](https://www.cvedetails.com/product/47/Linux-Linux-Kernel.html?vendor_id=33)

which allow this guarantee. The primary one is the unique system of ownership combined with the borrow checker. This is part of the compiler that ensures pieces of data live at least as long as they need to in order to be alive when they are used.

One other feature is the built-in `Option` type. This is used to replace the concept of null found in many other languages. In some languages, every type is secretly the union of that type with null. This means that you can always end up with bugs where you assume some variable had a value and it actually was inhabited by the dreaded null. Rust disallows this by not having null and instead having a type which can explicitly wrap other types. For example, consider this Rust code:

```
fn print_number(num: Option<i32>) {
    match num {
        Some(n) => println!("I see {}", n),
        None => println!("I see nothing!"),
    }
}

fn main() {
    let x = Some(42);
    let y = None;

    print_number(x);
    print_number(y);
}
```

The function `print_number` must handle the case where `num` is `None`, meaning the `Option` has no value. There are a few different ways to handle that case but you must explicitly do something for that case or else your code will not compile.

The one caveat here is that Rust does allow blocks of code to be marked `unsafe` and within those blocks it is possible to violate memory safety. Some things are impossible for the compiler to verify are safe and therefore it refuses to do so. It requires you to use `unsafe` regions of code to ensure that you understand the invariants required to make sure your code truly is safe.

This does not defeat the purpose, rather it isolates the areas of auditability to just those sections of code which are specifically marked. Nothing you do in normal Rust,

also called safe Rust, can result in a memory safety violation, unless something in unsafe code did something wrong ahead of you.

As an example, calling C functions from Rust is unsafe. This is because Rust has no way of knowing what the C code is doing, and C is inherently unsafe, therefore the compiler cannot uphold its guarantees if you call out to C. However, can it be safe to call C? Yes, provided you fill in the visibility gap for the compiler with your own logic.

## **Fearless concurrency**

Concurrency in programming means that multiple tasks can be worked on at the same time. This is possible even for a single thread of execution by interleaving the work for different tasks in chunks rather than only working on tasks as entire chunks.

Parallelism in programming means multiple tasks executing at the exact same time. True parallelism requires multiple threads of execution (or the equivalent).

The Rust language describes its facilities for concurrent and parallel computing as fearless concurrency with a bit of conflation of terms. I will continue in this tradition and use concurrency to mean concurrency and/or parallelism.

Most modern, high level languages have chosen how they want to support concurrency and mostly force you down that path. Some more general purpose languages provide the tools to handle concurrency however you see fit. For example, Go is designed around Communicating Sequential Processes (CSP) and therefore concurrency is most easily achieved using channels and goroutines. Python, on the other hand, has libraries for threads, multiprocesses, message passing actors, etc.

Rust is a low-level language by design and therefore provides tools that allow you to use the model of your choice to achieve your particular goals. Therefore, there are facilities for threads but also channels and message passing.

Regardless of what technique you choose to tackle concurrency and/or parallelism, the same ownership model and type system that ensures memory safety also ensures thread safety. This means that it is a compile time error to write to the same memory from different threads without some form of synchronization. The details are less important than the concept that entire classes of problems that are notoriously difficult to debug in other languages are completely eliminated at compile time while, importantly, retaining all of the performance benefits.

## C interoperability

Rust is foremost a systems programming language. That means it is designed for building low level systems with strict performance requirements and reliability constraints. Frequently in this world, C is the glue that binds many disparate systems. Therefore being able to interoperate with C is an absolute necessity to be able to have a serious systems language. Luckily it is straightforward to interact with C both by calling into C from Rust, as well as exposing Rust as a C library.

You might be saying that sounds great but I don't plan on writing an operating system anytime soon so why should I care? C is also the most common mechanism for making dynamic languages faster. Typically, when parts of your Python or Ruby code are showing performance problems, you can reach for an extension written in C to speed things up. Well, now you can write that extension in Rust and get all of the high level benefits of Rust and still make your Python or Ruby code think it is talking to C. This is also quite an interesting area for interacting with the JVM.

## Ecosystem

Software is not constructed in a vacuum, the practice of programming is often a community driven endeavor. Every language has a community whether it actively cultivates it or not. The ecosystem around a language includes the community of people, but also the tooling or lack thereof.

Rust has grown quite a lot in its short life and has gone through some growing pains as a result. However, the community has always been very welcoming and importantly the culture is a first-class citizen. Rust specifically has a community team as part of the governance structure of the language. This goes a long way to helping the language and ecosystem grow and mature.

We will cover much of the useful tooling that exists around Rust in detail below. However, suffice it to say that the tooling around the language is some of the best that exists. There have been a lot of learnings over the past twenty years about how to manage toolchains and dependencies and Rust has incorporated all of this quite well.

## The nature of programming

The systems and applications we are building today are different than 50 years ago, they are even different than 10 years ago. Therefore, it should not be too much of a stretch to say that the tools we use should also be different.

There is an explosion of embedded systems due to what is commonly called the Internet of Things. However, is C still the best tool for that job? Mission critical software that controls real objects that could lead to serious consequences in the case of failure should be using the best tool for the job. Rust is a serious contender in this space. For example, it is easy to turn off dynamic memory allocation while still being able to use a lot of the nice parts of the language.

The other explosion is continuing on the web. We have been in a web revolution for quite a while now, but things have not slowed down. The deficits of JavaScript are well known and have been addressed along quite a few paths. We have many languages which compile to JavaScript but provide nice features like type systems or a functional paradigm. However, there are fundamental performance and security issues with JavaScript regardless of how you generate it. WebAssembly (WASM) is a step in a different direction where we can compile languages like Rust to a format natively executable in the browser.

## Fun

Rust is fun to write. You will disagree with this and think I am crazy at some point while you are learning Rust. There is a learning curve which can be distinctly not fun. However, once your mental model starts to shift, you will find yourself having moments of pure joy when your code just works after the compiler gives you the okay.

## Why not Rust

Rust is just another programming language and as such is just another software project. This means it has built up some legacy, it has some hairy parts, and it has some future plans which may or may not ever happen. Some of this means that for any given project, Rust might not be the right tool for the job.

One area in which Rust might not be right is when interfacing with large C++ codebases. It is possible to have C++ talk to C and then have C talk to Rust and vice versa. That is the approach you should take today if possible. However, Rust does not have a stable ABI nor a stable memory model. Hence, it is not directly compatible with C++. You can incrementally replace parts of a system with Rust and you can build new parts in Rust, but plug-and-play interoperability with C++ is not a solved problem.

Furthermore, Rust takes time to learn. Now this is often cited as a reason for sticking with some other language because one is deemed an expert in that language. However, a counter point might be that you are not as much of an expert in that language as you might believe. A further counter point is that it might not matter, the other language might be fundamentally flawed enough that being an expert is irrelevant. Nonetheless, there are times where using the tool you know is the right answer.

The gap between learning Rust and knowing it from using it in anger is a bit bigger than in some other languages. Therefore the learning curve might seem steeper than you are used to. However, this is primarily because what is safe in Rust with the borrow checker helping you can be insane in other languages.

Type systems are amazing. You tell the computer some facts about your problem domain and it continually checks that those things are true and lets you know if you screw up. Yet there are valid programs which are inexpressible in a static type system. This is both theoretically true and actually happens in practice. Moreover, dynamic languages can frequently be more productive for small, isolated tasks. Sometimes the cost of the type system is not worth it.

## This book's mission

Rust has a great set of documentation around the standard library and has an official “book”<sup>5</sup> which is a great place to start if you are looking for another source of material. However, this book has a different focus than a traditional book trying to teach you a language. Our goal is to build realistic applications and explore some of the techniques and tools available in Rust for accomplishing those tasks.

---

<sup>5</sup><https://doc.rust-lang.org/book/>

In the process of working through some common scenarios, hopefully you will also be able to learn Rust. There is a gradual ramp up from very simple to more complex programs as we build up our Rust toolbelt. One specific goal is to show places where many people usually stumble and try to support you in finding your own ways over those hurdles. This should empower you when you branch out to your own problems.

This approach has the downside of not necessarily covering every language feature in the same depth or in the same order that you might encounter in a standard programming language introduction. Furthermore, we will explicitly try to take a pragmatic path rather than belabor esoteric details. Those details can be quite interesting and will be there for you when you want to seek them out, but often they get in the way of learning. We will sometimes do things in a less than perfect way as the trade-off is worth the expositional benefit.

Overall the goal is to get you to a state of productivity as quickly as possible. Along the way we will provide pointers to further material if you want to go deeper.

## **Setting expectations based on your background**

A great deal of terminology in the programming language space is built on a base set of shared ideas that have become so entrenched as to be overlooked by most every day developers. There are some lines we have to draw where we lean on some assumed prior knowledge. Thus, if you have never written any code before this might be a challenging book to make it entirely through. If you are willing to take some leaps of faith then you should be able to make it.

First and foremost, absolutely zero Rust background is assumed and every new concept will be explained as it arises.

If you have a background that only includes garbage collected, dynamically typed languages, such as Python, JavaScript, Ruby, PHP, then the biggest hurdle will probably be working with the type system. That being said, you might just find a wondrous joy associated with the tight feedback loop of a compiler telling you all the things you have done wrong. Moreover, the type inference will let you forget about it most of the time. Some of the points around memory safety might seem less exciting to you because all of these languages are also memory safe. The approach



to achieving memory safety is different but the end result is the same. Some topics around pointers and references might seem new, but all of these languages leak those concepts and you probably already understand them just in a different form. For example, if you write the following Python:

```
def someFunc(items = []):  
    items.append(1)  
    return items
```

```
a = someFunc()  
b = someFunc()
```

```
a.append(2)
```

```
print(a)  
print(b)
```

you will see `[1, 1, 2]` printed twice. As a diligent Python programmer you know not to use lists as default arguments to functions like this and the reason has to do with values versus references. So even if you don't explicitly think that you are working with pointers, you definitely do use them all the time. Rust has many high level syntactic features that make it feel surprisingly similar to a dynamic language.

If you have a background in functional programming coming from Haskell or the ML family then a lot will feel quite at home. But, the use of mutability and explicit imperative style might be less of your thing. Rust has great functional programming facilities and the type system borrows a lot from these languages. However, Rust is more geared towards giving you a lot of the same safety benefits of functional programming while still writing imperative code. Shared, mutable state is the root of all evil. Functional programming attacks that problem by doing away with mutability. Rust attacks it by doing away with sharing.

If you are coming from C++ then you are in for an easier road in some ways and a much harder one in others. I focus here on C++ as it is more superficially similar, but many points also apply to C. Much of the syntax and most of the concepts will be familiar to you. However, there are a few new concepts, like lifetimes, and a few things that look the same but are not, like references and move semantics.

There are APIs you will find in Rust which you might find to be highly performant, but laughably dangerous if ported to C++. You are correct. However, the borrow checker can make such APIs safe. For example, would you give a reference to a stack allocated piece of data to another thread? Would you store a reference to part of a string in a heap allocated struct? Both those are invitations to disaster in C++. They are trivial to do correctly in Rust thanks to the borrow checker and can be great for performance. The API you might find normal in C++ may not be expressible in safe Rust. That is, the borrow checker may not allow you to compile code you consider correct. You may in fact be correct. However, this is usually an API which is easy to misuse and is only correct with a significant amount of cognitive burden on you.

Hence, coming from C++ might require the most shift in how you think about structuring programs. You are more likely to “fight the borrow checker” because some of the ways Rust wants you to do things are just plain against your instincts.

Rust has a bit of notoriety for having a steep learning curve, but it is actually mostly about unlearning things from other languages. Therefore, having less experience can work in your favor.

## Getting your environment setup

This book assumes an installed version of Rust and some associated tooling. The first step in getting setup is to visit the official installation website:

<https://www.rust-lang.org/tools/install>

You should be able to follow the instructions to get setup via `rustup`. Rust has a fast release cadence for a programming language with a new version every six weeks. This means that the particular version as of this writing will be stale by the time you are reading it. However, Rust also puts a strong emphasis on backwards compatibility. Thus, as long as you are using a version of Rust at least as new as when this was written, everything should still work for you. Rust 1.37.0 should be new enough for all the code in this book. Moreover, we are using the 2018 edition exclusively. There is an [entire guide](#)<sup>6</sup> dedicated to explaining the editions so we will not cover it in depth here.

---

<sup>6</sup><https://doc.rust-lang.org/edition-guide/index.html>

## Rustup

The [rustup](https://rustup.rs)<sup>7</sup> tool is your one stop shop for managing multiple versions of the Rust compiler on your machine. You can have different versions of the compiler installed next to each other and easily switch back and forth between them. You can install nightly releases to try out new features and then easily switch back to stable for other projects. If you have ever dealt with the absolute madness associated with managing different versions of some languages then you will be delighted at how well rustup just works.

One note, for some reason all of the details of rustup can be found in [the Github readme](https://github.com/rust-lang/rustup)<sup>8</sup> for the project. It is pretty easy to use but the command line help frequently fails me.

## Cargo

rustc is the Rust compiler, and you can invoke it directly, however you will find this rarely to be necessary as the majority of your time will be spent interacting with Cargo. Cargo is a dependency manager and a build system. You use a manifest to specify details of your code and its dependencies and you can then instruct Cargo to build your code and it will take care of the rest. You can have Cargo manage building for other platforms and for quickly type checking via `cargo check`. You use it to run tests via `cargo test` and for countless other tasks.

We will cover code structure later on, but the primary unit is known as a crate. You can depend on other crates and the public repository can be found at [crates.io](https://crates.io)<sup>9</sup>. This is related to Cargo in that there is quite a bit of default work built in to Cargo for working with crates.io, but it is not absolutely required.

Cargo has [its own guide](https://doc.rust-lang.org/cargo/guide/)<sup>10</sup> which is a great source of information when you find yourself wondering how to do something with Cargo. You can also always run `cargo help` to answer your questions from the command line.

---

<sup>7</sup><https://rustup.rs>

<sup>8</sup><https://github.com/rust-lang/rustup>

<sup>9</sup><https://crates.io/>

<sup>10</sup><https://doc.rust-lang.org/cargo/guide/>

## IDEs, RLS, Editors

The editor support story is getting better and is significantly better than it used to be. A project known as the [Rust Language Server](https://github.com/rust-lang/rls)<sup>11</sup>(RLS) is designed to provide the backend for any editor to interact with the compiler and a tool called [Racer](https://github.com/racer-rust/racer)<sup>12</sup> which provides faster (but less precise) information than the compiler can. This is a project that conforms to the [Language Server Protocol](https://langserver.org/)<sup>13</sup>(LSP) so that every editor which can act as a LSP client can work with RLS. There is a reference implementation of an RLS specific frontend for Visual Studio Code, so if you are unsure where to start that might be one to try out.

If you have a favorite editor already, like Vim or Emacs, then there are plugins you can use to make working with Rust more comfortable. Personally, I use Vim and a shell for running commands directly with Cargo. This is mostly so that I can move between environments with minimal change to my workflow, and I have found that Rust is amenable to this style. There are some languages which are very hard to work with without autocomplete and Rust has not been like that for me.

Check out the [official website](https://www.rust-lang.org/tools)<sup>14</sup> for an up to date list of tools.

## Clippy

The linter is affectionately named [Clippy](https://github.com/rust-lang/rust-clippy)<sup>15</sup>. Cargo supports an awesome feature where you can install subcommands via `rustup` so that you can selectively add components to Cargo based on your needs. Clippy can be installed this way by running:

```
rustup component add clippy
```

and then run with:

---

<sup>11</sup><https://github.com/rust-lang/rls>

<sup>12</sup><https://github.com/racer-rust/racer>

<sup>13</sup><https://langserver.org/>

<sup>14</sup><https://www.rust-lang.org/tools>

<sup>15</sup><https://github.com/rust-lang/rust-clippy>

```
cargo clippy
```

It provides a bunch of helpful information and is good to run against your code regularly. There are many ways to configure it both at a project level as well as at particular points in your code. Linters are still an under used tool that end up being a big source of bike shedding on larger teams. However using at least something to catch the most egregious issues is better than nothing.

## Rustfmt

Rust has an official code formatter called `rustfmt`<sup>16</sup>. This was a project that started life in the community and eventually got official status. However, it is not as seriously official as `gofmt` for the Go language. You can configure `rustfmt` based on a couple attributes and there is nothing forcing you to use it. But, you should use it. Automated code formatting is one of the great productivity wins of the past twenty years. Countless engineering hours will no longer be wasted debating the finer points of column widths, tabs versus spaces, etc. Let the formatter do its job and get back to building.

## Documentation

The standard library has `documentation`<sup>17</sup> which you will consult frequently. It is thorough and well-written. I know that typing `d` into my browser's address bar and hitting enter will take me to `doc.rust-lang.org`.

All crates on `crates.io`<sup>18</sup> will automatically have its documentation built and available on `docs.rs`<sup>19</sup> which is an amazing tool for the community. Rust has great facilities for including documentation in your code which is why most crates are quite well documented. One excellent feature is the ability to include code samples in your documentation which is actually checked by the compiler. Thus the code examples in the documentation are never out of date

---

<sup>16</sup><https://github.com/rust-lang/rustfmt>

<sup>17</sup><https://doc.rust-lang.org/std/index.html>

<sup>18</sup><https://crates.io>

<sup>19</sup><https://docs.rs/>

The [official rustdoc book](https://doc.rust-lang.org/rustdoc/index.html)<sup>20</sup> is a great resource for learning about documenting your Rust code.

## The Nomicon

Rust has a safe and an unsafe side. You may one day find yourself wondering more about what goes on over on the mysterious, dark unsafe side. Well look no further than the [Rustonomicon](https://doc.rust-lang.org/nomicon/)<sup>21</sup> known colloquially as The Nomicon. This book will give you more guidance about how safe and unsafe interact, what you can and cannot do in unsafe Rust, and most importantly how not to break things when you need to use unsafe. It is highly unusual to need to use unsafe. Even when performance is critical, safe Rust most likely can solve your problem. However, there are instances where you really need to reach for this tool and the Nomicon will be your friend at that time.

---

<sup>20</sup><https://doc.rust-lang.org/rustdoc/index.html>

<sup>21</sup><https://doc.rust-lang.org/nomicon/>

# Summary

The history of computing is filled with powerful abstractions that let the machine manage complexity and thus frees cognitive load from our minds to be spent on more productive tasks. We moved from machine code to assembly to high level languages like C. Each step had a cost associated with giving up some explicit control and a benefit of increased expressive power per unit of code written. Some of these layers were strictly better, i.e. the cost was so negligible compared to the benefits as to be ignored today.

These layers of abstractions continued to be built with different cost/benefit trade-offs. Writing programs to solve complex tasks is a challenging endeavor. Some constraints on our programs have dictated which of those layers of abstractions we can use based on the various trade-offs. At some point, it became generally accepted wisdom that memory safety must be traded off against performance and control. If you want performance then you have to be close to the metal and that is necessarily unsafe. If you want safety then you must be willing to sacrifice some runtime performance to get it. There were counter points in small programs, but no one has challenged this status quo when applied to programming in the large. That is, until Rust.

You can have speed, no garbage collection and therefore a low memory footprint, and you can have safety. Rust affirms our worst fears: programming is hard, humans are fallible. But Rust also assuages those fears by having our computer handle the tasks that are hard for humans. The concepts behind Rust are just the natural evolution in our history of layered abstractions. Whether Rust succeeds or not, history will look back at this as a turning point where it no longer became acceptable to give up safety.

# Making Your First Rust App

## Getting started

We are going to build an application in Rust to get a feel for the language and ecosystem. The first step for all new Rust projects is generating a new project. Let's create a new project called `numbers`:

```
cargo new numbers
```

Cargo is the package manager for Rust. It is used as a command line tool to manage dependencies, compile your code, and make packages for distribution. Running `cargo new project_name` by default is equivalent to `cargo new project_name --bin` which generates a *binary* project. Alternatively, we could have run `cargo new project_name --lib` to generate a *library* project.

## Binary vs. library

A binary project is one which compiles into an executable file. For binary projects, you can execute `cargo run` at the root of your application to compile and run the executable.

A library project is one which compiles into an artifact which is shareable and can be used as a dependency in other projects. Running `cargo run` in a library project will produce an error as cargo cannot figure out what executable you want it to run (because one does not exist). Instead, you would run `cargo build` to build the library.

There are different formats which the Rust compiler can generate based on your configuration settings depending on how you wish to use your library.

The default is to generate an `rlib` which is a format for use in other Rust projects. This allows your library to have a reduced size for further distribution to other Rust



projects while still being able to rely on the standard library and maintain enough information to allow the Rust compiler to type check and link to your code.

Alternative library formats exist for more specialized purposes. For example, the `cdylib` format is useful for when you want to produce a dynamic library which can be linked with C code. This produces a `.so`, `.dylib`, or `.dll` depending on the target architecture you build for.

## The generated project

Let's enter the directory for our newly generated Rust project to see what is created:

```
cd numbers
```

The generated structure is:

```
.
├── Cargo.toml
└── src
    └── main.rs
```

Rust code organization relies primarily on convention which can be overridden via configuration, but for most use cases the conventions are what you want.

### main.rs

For a binary project, the entry point is assumed to be located at `src/main.rs`. Furthermore, inside that file, the Rust compiler looks for a function named `main` which will be executed when the binary is run. Cargo has generated a `main.rs` file which contains a simple “Hello, world!” application:

src/main.rs

```
1 fn main() {  
2     println!("Hello, world!");  
3 }
```

The syntax here says define a function (`fn`) with the name `main` which takes zero arguments and returns the empty tuple `()`.

Leaving off the return type is equivalent to writing `-> ()` after the argument list of the function. All function calls are expressions which must return a value. The empty tuple `()` is a marker for no value, which is what a function with no return type implicitly returns.

The body of the function calls a macro `println` which prints its argument `"Hello, world!"` to standard out followed by a newline.

We will cover macros more later, but for now we will mention a few basics. We know it is a macro invocation and not a normal function call because of the trailing `!` in the name. Macros are a powerful form of meta-programming in Rust which you will use frequently but probably rarely find the occasion to have to write. Rust implements `println` as a macro instead of as a regular function because macros can take a variable number of arguments, while a regular function cannot.

The syntax of Rust is superficially similar to C++ which we can see as curly braces are used for denoting blocks and statements are semicolon terminated. However, there is quite a bit more to the Rust grammar that we will cover as we go along.

## Cargo.toml

The `Cargo.toml` file is the manifest file for the project which uses the TOML<sup>22</sup> format. This is the entry point for describing your project as well as specifying

---

<sup>22</sup><https://github.com/toml-lang/toml>

dependencies and configuration. The initial generated file contains the bare essentials for describing your project:

Cargo.toml

---

```
1 [package]
2 name = "numbers"
3 version = "0.1.0"
4 authors = ["Your Name <your.name@example.com>"]
5 edition = "2018"
6
7 [dependencies]
```

---

The blank section for dependencies is included because nearly every project includes some dependencies. One feature of Rust has been to keep the core language and standard library relatively slim and defer a lot of extra functionality to the community. Therefore relying on third party dependencies is encouraged.

## Crates

The primary unit of code organization in Rust is called a crate. Your code exists as a crate which can be distributed to the community via [crates.io](https://crates.io/)<sup>23</sup>. Crates in Rust are analogous to gems in Ruby or packages in JavaScript. The registry at crates.io is similar to rubygems.org or npmjs.com as the de facto community repository for distributing and sharing code.

Binary Rust projects are also called crates so they do not solely represent shared library code. Furthermore, a crate can contain both a library and an executable.

It is often difficult to foresee how other's will want to use your software. A common practice in the Rust community is to create dual library/binary crates even when the primary intention of a project is to produce an executable. This can have positive effects on the API design of your code knowing that it should be suitable for external consumption. The binary part of the crate is typically responsible for argument parsing and configuration, and then calls into the functionality exposed by the library

---

<sup>23</sup><https://crates.io/>

part of the crate. Writing all of your code only as an executable and then trying to extract a library after the fact can be a more difficult process. Moreover, the cost of splitting code into a library is minimal.

## Making our crate a library

Cargo assumes the entry point for defining a library crate is a file `src/lib.rs`. Let's convert our current binary crate into a binary and library crate. First, we create our library entry point:

`src/lib.rs`

---

```
1 pub fn say_hello() {  
2     println!("Hello, world!");  
3 }
```

---

There are two differences to this code from what was in `main.rs`. First, we changed the name of the function from `main` to `say_hello`. This change is more cosmetic than anything (in fact leaving it named `main` works just fine, `main` is only special in some contexts).

The second change is the keyword `pub` before `fn`. This is a privacy identifier which specifies that this function should be publicly accessible to users of our crate. Without the keyword, we could call this function inside of our `lib.rs` file, but users of our crate would not be able to call it. Note that our executable sees the library crate the exact same as someone who included our library as a dependency in their `Cargo.toml` file. This ensures a proper separation of concerns between code meant to be executed as a binary and the actual functionality of your project.

We can now change our `main` function to use the functionality exposed by our library:

src/main.rs

```
1 fn main() {  
2     numbers::say_hello();  
3 }
```

Running this code should result in the same output as before:

```
$ cargo run  
    Compiling numbers v0.1.0 (...)  
    Finished dev [unoptimized + debuginfo] target(s) in 0.53s  
    Running `target/debug/numbers`  
Hello, world!
```

Let's unpack this function call syntax a little bit before moving on. Even though our binary exists in the same codebase as our library, we still must refer to the functions in the crate by the name of the crate, `numbers` in this case.

We wish to call a function named `say_hello` which exists in the `numbers` crate. The double colon operator `::` is used for separating items in the hierarchy of modules. We will cover modules later, but suffice it to say that crates can contain modules, which themselves can contain more modules.

To resolve an item, be it a type or function, you start with the name of the crate, followed by the module path to get to the item, and finally the name of the item. Each part of this path is separated by `::`. For example, to get a handle to the current thread you can call the function `std::thread::current`. The crate here is `std` which is the standard library. Then there is a module called `thread`. Finally inside the `thread` module there is an exported function called `current`.

Items can exist at the top level of a crate (i.e. not nested in any modules), which you refer to simply by the name of the crate, then `::`, then the name of the item. This is what is happening with `numbers::say_hello` because `say_hello` exists at the top level of our `numbers` crate.

## Trade-offs

Two of the big selling points of Rust are performance and reliability. Performance meaning both runtime speed and memory consumption. Reliability here means catching bugs at compile time and preventing certain classes of errors entirely through language design. These goals are often seen as classically at odds with one another. For example, C lives in a world where performance is of utmost importance, but reliability is left as an exercise for the implementor.

Rust has no garbage collector and no runtime in the traditional sense. However, most difficulties of working with manual memory management are taken care of for you by the compiler. Therefore, you will often hear “zero cost” being used to describe certain features or abstractions in the language and standard library. This is meant to imply that neither performance nor reliability has to suffer to achieve a particular goal. You write high level code and the compiler turns it into the same thing as the “best” low level implementation.

However, in practice, what Rust really gives you is the tools to make choices about what trade-offs you want to make. Underlying the design and construction of all software is a series of trade-offs made explicitly or implicitly. Rust forces more of these trade-offs to be made explicitly which can make the initial learning period seem a bit more daunting especially if you have experience in languages where many trade-offs are implicit.

This will be a topic that permeates this book, but for now we will highlight some of these aspects as we make our numbers crate do something more interesting.

## Print a list of numbers

Let’s build an application that creates a list of numbers and then prints each number on a line by itself to standard out. As a first step, let’s just say we want to print the numbers one through five. Therefore, our goal is the following:

```
$ cargo run
1
2
3
4
5
```

Let's change our `main` function to call the yet to be defined library function `print`:

**src/main.rs**

---

```
1 fn main() {
2     numbers::print();
3 }
```

---

Since we want to print one through five, we can create an array with those numbers and then print them out by looping over that array. Let's create the function `print` in `lib.rs` to do that:

**src/lib.rs**

---

```
1 pub fn print() {
2     let numbers = [1, 2, 3, 4, 5];
3     for n in numbers.iter() {
4         println!("{}", n);
5     }
6 }
```

---

Let's unpack this from the inside out. We have already seen the `println` macro, but here we are using it with a formatted string for the first time. There are two main features of string interpolation in Rust that will take you through most of what you need. The first argument to one of the printing macros (`print`, `println`, `eprint`, `eprintln`) is a double quoted string which can contain placeholders for variables. The syntax for placeholders to be printed “nicely” is `{}`, and for debugging purposes is `{:?}`. The full syntax for these format strings can be found [in the official docs](https://doc.rust-lang.org/std/fmt/)<sup>24</sup>.

---

<sup>24</sup><https://doc.rust-lang.org/std/fmt/>

The “nice” format is possible when a type implements the `Display` trait. The debugging format is possible when a type implements the `Debug` trait. Not all types implement `Display`, but the standard practice is for all public types to implement `Debug`. So when in doubt, use `{:?}` to see the value of some variable and it should almost always work.

We will cover traits in detail later, but we will give a crash course here for what is necessary. Traits are part of the type system to mark certain facts about other types. Commonly they are used to define an interface to a particular set of functions that the type in question implements. You can define your own traits as well as implement traits. Whether a type implements a trait must be stated explicitly in code rather than implicitly by satisfying the functional requirements of the trait. This is one of a few differences between Rust traits and Go interfaces.

We will see when creating types later that usually you can get a `Debug` implementation derived for free, but you must implement `Display` yourself if you want it. Most built-in types implement `Display`, including integers, so we use the format string `"{}"` to say expect one variable. Note that the following does not work:

```
println!(n);
```

The first argument to the `print` macros must be a literal string, it cannot be a variable, even if that variable points to a literal string. Therefore, to print out a variable you need to use the format string `"{}"` as we are doing. If you forget this the Rust compiler will suggest that as what you probably want to do.

## Iteration

So assuming that `n` holds an integer from our collection, we are printing it out using the `println` macro. How does `n` get bound to the values from our collection? We loop over our collection using a `for` loop and bind `n` to each value. The syntax of a `for` loop is:



```
for variable in iterator {  
    ...  
}
```

Note that we are calling the method `iter` on our array. Rust abstracts the idea of iteration into yet another trait, this one called `Iterator`. We have to call `iter` here to turn an array into an `Iterator` because arrays do not automatically coerce into an `Iterator`. We shall see shortly that this is not always necessary with other collections.

This is also the first time we are calling a method on an object. Rust types can implement functions that operate on themselves and can therefore be called using this dot syntax. This is syntactic sugar for a direct function call with the receiver object as the first argument. We will cover how these functions are defined when we construct our own types and implement methods on them.

## Defining Array Types

We can move out further now to the definition of our array. Rust borrows many ideas of the ML family of languages so some concepts might be familiar if you have experience in that area. By default variables are immutable. Therefore we declare an immutable variable called `numbers` which is bound to an array with the numbers we are interested in. Rust infers the type of `numbers` based on the value we used to initialize the variable. If you want to see the type that is inferred by Rust, a trick is to write:

```
let () = numbers;
```

after the line that declares the variable `numbers`. When you try to compile this code, there will be a type mismatch in the assignment which will print out what the compiler expects:

```
$ cargo run
  Compiling numbers v0.1.0 (...)
error[E0308]: mismatched types
  --> src/lib.rs:3:9
   |
3 |     let () = numbers;
   |           ^^ expected array of 5 elements, found ()
   |
   = note: expected type `[integer]; 5`
           found type `()`
```

```
error: aborting due to previous error
```

For more information about this error, try `rustc --explain E0308`.  
 error: Could not compile `numbers`.

To learn more, run the `command` again with `--verbose`.

We see that the compiler inferred a type of `[integer; 5]` for `numbers`. Arrays in Rust are a homogeneous container (all elements have the same type) with a fixed size. This allows it to be stack allocated. The ability to ensure data is stack allocated rather than heap allocated is one of the areas in which Rust allows you to decide what trade-offs you want to make. On the other hand, because an array has a fixed size that must be known at compile time it is not useful for data which might need to grow or shrink or contain an unknown numbers of items. For this we have the `Vec` type which we will return to shortly.

You can also see that the compiler infers the type of elements of the array to be `{integer}` which is a placeholder as without any further constraints the specific type of integer is unknown. Rust has twelve integer types which depend on size and whether it is signed or unsigned. The default is `i32` which means a signed integer that takes 32 bits of space. The equivalent unsigned type is `u32`. Let's say we wish our numbers to be `u8`, that is 8-bit unsigned integers. One way to do this is to specify directly on the numerical constant what type we want:

```
let numbers = [1u8, 2, 3, 4, 5];
```

If we do this then the compiler will infer the type `[u8; 5]` for our array. The other way is to explicitly write out the type of the variable `numbers`:

```
let numbers: [u8; 5] = [1, 2, 3, 4, 5];
```

Type annotations are written with a colon after the variable name followed by the type. We see that the size of the array (5) is part of the type. Therefore, even with the same type of elements, say `u8`, an array with four elements is a different type than an array with five elements.

## Using `std::vec::Vec`

Rust provides a few mechanisms for alleviating some of the limitations of arrays. The first we will talk about is the vector type in the standard library, `std::vec::Vec`<sup>25</sup>. A vector is similar to an array in that it stores a single type of element in a contiguous memory block. However, the memory used by a vector is heap allocated and can therefore grow and shrink at runtime. Let's convert our library print function to use a vector:

`src/lib.rs`

---

```
1 pub fn print() {  
2     let numbers = vec![1, 2, 3, 4, 5];  
3     for n in numbers {  
4         println!("{}", n);  
5     }  
6 }
```

---

We are calling another macro `vec` which this time constructs a vector with the given values. This looks very similar to the array version, but is actually quite different. Vectors own their data elements, they have a length which says how many elements are in the container, and they also have a capacity which could be larger than the

---

<sup>25</sup><https://doc.rust-lang.org/std/vec/struct.Vec.html>

length. Changing the capacity can involve quite a bit of work to allocate a new region of memory and move all of the data into that region. Therefore, as you add elements to a vector, the capacity grows by a multiplicative factor to reduce how frequently this process needs to take place. The biggest advantage is that you do not need to know upfront how large the vector needs to be; the length is not part of the type.

The type of a vector is `Vec<T>` where `T` is a generic type that represents the types of the elements. Therefore, `Vec<i32>` and `Vec<u8>` are different types, but a `Vec<u8>` with four elements is the same type as one with five elements.

Note also that we are no longer explicitly calling `iter` on the `numbers` variable in our `for` loop preamble. The reason for this is that `Vec` implements a trait that tells the compiler how to convert it into an iterator in places where that is necessary like in a `for` loop. Calling `iter` explicitly would not be an error and would lead to the same running code, but this implicit conversion to an iterator is common in Rust code.

## Function Arguments

Let's abstract our `print` function into two functions. The entry point will still be `print` (so we don't need to change `main`) which will construct a collection, but it will then use a helper function to actually print the contents of this collection. For now we will go back to using an array for the collection:

`src/lib.rs`

---

```
1 pub fn print() {
2     let numbers = [1, 2, 3, 4, 5];
3     output_sequence(numbers);
4 }
5
6 fn output_sequence(numbers: [u8; 5]) {
7     for n in numbers.iter() {
8         println!("{}", n);
9     }
10 }
```

---

This is our first function that has input or output. Type inference does not operate on function signatures so you must fully specify the types of all inputs and the output.

However, we still are not returning anything so by convention we elide the `-> ()` return type which is the one exception to the rule of fully specifying the types in function signatures.

The input type of our function `output_sequence` is our five element array of `u8` values.

Rust has a few different modes of passing arguments to functions. The biggest distinction being that **Rust differentiates between:**

- a function temporarily having access to a variable (borrowing) and
- having *ownership* of a variable.

Another dimension is whether the function can mutate the input.

The default behavior is for a function to take **input by value and hence ownership** of the variable is moved into the function.

The exception to this rule being if the type implements a special trait called `Copy`, in which case the input is copied into the function and therefore the caller still maintains ownership of the variable. If the element type of an array implements the `Copy` trait, then the array type also implements the `Copy` trait.

Suppose we want to use a vector inside `print` instead, so we change the code to:

`src/lib.rs`

---

```
1 pub fn print() {  
2     let numbers = vec![1, 2, 3, 4, 5];  
3     output_sequence(numbers);  
4 }  
5  
6 fn output_sequence(numbers: [u8; 5]) {  
7     for n in numbers.iter() {  
8         println!("{}", n);  
9     }  
10 }
```

---

But this won't work because `[u8; 5]` and `Vec<u8>` are two different types. One possible fix is to change the input type to `Vec<u8>`:

src/lib.rs

---

```
1 pub fn print() {
2     let numbers = vec![1, 2, 3, 4, 5];
3     output_sequence(numbers);
4 }
5
6 fn output_sequence(numbers: Vec<u8>) {
7     for n in numbers {
8         println!("{}", n);
9     }
10 }
```

---

This works for this case. It also let's us see what happens when passing a non-Copy type to a function. While arrays implement the Copy trait if their elements do, Vec does not. Hence, try adding another call to `output_sequence(numbers)` after the first one:

src/lib.rs

---

```
1 pub fn print() {
2     let numbers = vec![1, 2, 3, 4, 5];
3     output_sequence(numbers);
4     output_sequence(numbers);
5 }
6
7 fn output_sequence(numbers: Vec<u8>) {
8     for n in numbers {
9         println!("{}", n);
10    }
11 }
```

---

This gives us an error:

```
$ cargo run
   Compiling numbers v0.1.0 (...)
error[E0382]: use of moved value: `numbers`
  --> src/lib.rs:4:21
   |
3 |     output_sequence(numbers);
   |                      ----- value moved here
4 |     output_sequence(numbers);
   |                      ^^^^^^^^ value used here after move
   |
   = note: move occurs because `numbers` has type `std::vec::Vec<u8>`, which does not implement the `Copy` trait

error: aborting due to previous error

For more information about this error, try `rustc --explain E0382`.
error: Could not compile `numbers`.
```

To learn more, run the command again with `--verbose`.

We can see Rust generally has very helpful error messages. The error is that a value was used after it has been moved. The `print` function no longer owns `numbers`. The “note” in the error explains why the move happens due to vector not implementing the `Copy` trait.

Note that in the changes we have made, the body of `output_sequence` has remained the same (modulo whether we call `iter` explicitly or not), only the type signature has been changing. This is a hint that maybe there is a way to write a type signature that works for both arrays and vectors. There are again several ways to accomplish this goal.

## A type signature that works for both arrays and vectors

As we have said before, Rust has a lot of power and gives you very fine-grained control over what you want to use or don’t want to use. This can be frustrating when starting out because any time you ask “what is the right way to do this,” you

will invariably be met with the dreaded “it depends.” Rather than detail every possible permutation that achieves roughly the same outcome, we are going to focus on the most common idioms. There are certain performance reasons as well as API design decisions that lead to different choices, but those are more exceptional cases than the norm. We will provide pointers to the choices we are making when it matters, but note that due to the scope of the language there is almost always more than one way to do it.

A key type that comes in handy to alleviate some of the limitations of arrays is the `std::slice`<sup>26</sup>. Slices are a dynamically sized view into a sequence. Therefore, **you can have a slice which references an array or a vector and treat them the same**. This is a very common abstraction tool used in Rust. This will be more clear by seeing this in action.

Let’s change the signature of `output_sequence` to take a reference to a slice, and change `print` to show that it works with both arrays and vectors:

src/lib.rs

---

```
1 pub fn print() {
2     let vector_numbers = vec![1, 2, 3, 4, 5];
3     output_sequence(&vector_numbers);
4     let array_numbers = [1, 2, 3, 4, 5];
5     output_sequence(&array_numbers);
6 }
7
8 fn output_sequence(numbers: &[u8]) {
9     for n in numbers {
10         println!("{}", n);
11     }
12 }
```

---

A slice of `u8` values has type `[u8]`. This represents a type with an unknown size at compile time. The Rust compilation model does not allow functions to directly take arguments of an unknown size. In order to access this slice of unknown size with something of a known size we use indirection and pass a reference to the slice

---

<sup>26</sup><https://doc.rust-lang.org/std/slice/index.html>



rather than the slice itself. A reference to a slice of `u8` values has type `&[u8]` which has a known size at compile time. This size is known because it is equal to the size of a pointer plus the length of the slice. Note that slices convert automatically into iterators just like vectors so we again do not call `iter` explicitly in the body of our function. This takes care of the signature of `output_sequence` however the way we call this function from `print` has changed as well.

Notice that we have added an `&` before the variable names that are passed to `output_sequence`. You can think of this as creating a slice that represents read-only access to the entire sequence for both the vector and array. However, this small change in how we call the function allows us to handle vectors and arrays equally well. Idiomatic Rust takes slices as arguments in most cases where one needs only to read the collection. This is particularly true for strings which we will cover later.

The major difference here is that we are no longer transferring ownership into the function `output_sequence` instead we are lending read-only access to that function. The data is only borrowed for the duration of the function call. The idea of ownership and borrowing is a core part of the Rust language and is something we will be constantly running into.

## Constructing A Vector of Numbers

Let's make one more change to make this program more flexible. Instead of printing out one through five, let's take a number as input and print from one up to that value. We could just iterate through integers and print them out as we go along rather than using the `output_sequence` helper function. However, we are going to construct a vector to show a few more language features.

Let's create yet another helper function, `generate_sequence` which takes a limit as input and outputs a vector. Our `print` function can then just combine these two parts:

src/lib.rs

```
1 pub fn print(limit: u8) {  
2     let numbers = generate_sequence(limit);  
3     output_sequence(&numbers);  
4 }  
5  
6 fn generate_sequence(limit: u8) -> Vec<u8> {  
7     let mut numbers = Vec::new();  
8     for n in 1..=limit {  
9         numbers.push(n);  
10    }  
11    numbers  
12 }  
13  
14 fn output_sequence(numbers: &[u8]) {  
15     for n in numbers {  
16         println!("{}", n);  
17     }  
18 }
```

In `print` we bind a variable to the result of calling `generate_sequence` with the `limit` passed to us as the argument, then we call `output_sequence` as before passing a reference to a slice backed by the variable we just created.

The new function here takes an input argument, `limit`, and returns a `Vec<u8>`. This is our first function returning something.



Again as there are a lot of different ways to do things in Rust, we are going to just show one particular way to construct the vector we desire in order to hit some relevant parts of the language.

First we create a new vector with `Vec::new()`.



Unlike in some other languages, `new` is not special but rather has become by convention the name of the function that returns a new instance of a type. You can write a function called `new` which does something else and it would compile just fine, but it would go against the standard way of doing things.

By default a vector created with `new`, is the same as one created with `vec![]`, and does not allocate. Therefore, unless you actually put something into a vector it does not use any memory.

In the code above, we see a `new` keyword being used, `mut`. Mutability is a property of the variable or reference not of the object itself. Therefore we declare `numbers` to be a mutable variable that holds an empty vector. This allows us to later call `numbers.push(n)` because `push` is a method that requires the receiver to be mutable. Removing the `mut` from the `let` binding will result in a compiler error when we try to `push`.

In order to generate the numbers starting at 1 up to our limit, we use a `for` loop, but this time the iterator is a `Range`<sup>27</sup> object, in particular an `InclusiveRange`. Ranges can be constructed with using the syntax `start..end` or `start..=end`. Both `start` and `end` are optional, and if you have neither, i.e. `..`, then you also cannot have the `=` sign. By default the range is inclusive on the left (i.e. includes `start`), and exclusive on the right (i.e. does not include `end`). The `=` after the two dots makes it so the range includes the end point. We want the numbers starting at 1 up to `limit`, including the `limit`, so we use `1..=limit`. Ranges are frequently used when creating slices, for example:

```
let numbers = [1, 2, 3, 4, 5];
let subset = &numbers[1..3];
```

Here `subset` is a slice of length  $3-1=2$  which starts at index 1, hence it is the slice `[2, 3]`.

Iterating over this range, we push each value onto the end of our vector which causes heap allocations every time there is not enough capacity to extend the length. Finally, we want to return this vector from our function. The final expression in a function is

---

<sup>27</sup><https://doc.rust-lang.org/std/ops/struct.RangeInclusive.html>

implicitly returned so there is no need for an explicit `return` statement. However note the lack of semicolon at the end of the last line of this function. The expression that evaluates to the vector `numbers` is written without a semicolon and means to return that value. If we had written a semicolon, that would be a statement whose value is `()` which is not what you want to return. This is a common error so the compiler is smart enough to tell you what to fix, but it is nonetheless an error. You can use a `return` statement to return early from a function, but using the last expression of the block as the implicit return is idiomatic Rust.

## A Shorter Version with `collect`

The purpose of writing `generate_sequence` like this was to demonstrate object construction, mutability, and ranges. Before leaving, let's look at a very powerful construct that is used throughout real Rust which would be a better approach to generating this vector. We could replace `generate_sequence` with:

`src/lib.rs`

---

```
6 fn generate_sequence(limit: u8) -> Vec<u8> {  
7     (1..=limit).collect()  
8 }
```

---

Rust has powerful generic programming facilities which allows for the function `collect` to exist. This function can be used to turn any iterator into basically any collection.

Commonly, one takes a collection like a vector, turns it into an iterator by calling `iter`, performs transformations on the generic iterator, and then calls `collect` at the end to get back whatever specific collection one wants to work with. It can also be used to directly turn one collection into another collection, which is what we are doing here by turning a range into a vector.

`Collect` is a generic function over the return type, so the caller gets to determine what they want. Here because we return the result of calling `collect` from our function, type inference sees that the return type needs to be a `Vec<u8>` and therefore ensures that `collect` generates that collection. While the type inference in Rust is good, it some times cannot figure out what you want when using `collect`. Therefore, you

might find the need to use the syntax `collect::<SomeType>()` to help the compiler know what you want.



This syntax, `::<>`, you may see referred to as the “turbofish”.

We have changed our exported `print` function to require an input variable, so we need to update our call in `main` to pass something in. Let’s pass in 5 to get the same output as before:

`src/main.rs`

---

```
1 fn main() {  
2     numbers::print(5);  
3 }
```

---

## Testing our code

Testing is a large topic and is something we will cover in more detail as we move to larger applications, however let’s write our first test to see how easy it can be in Rust. Add the following to the end of `src/lib.rs`:

`src/lib.rs`

---

```
20 #[test]  
21 fn generate_sequence_should_work() {  
22     let result = generate_sequence(3);  
23     assert_eq!(result, &[1, 2, 3]);  
24 }
```

---

Our test is just a normal function with a special attribute, `#[test]`, before it. We will see attributes come up frequently as they are used for a variety of purposes in Rust. They come in two forms `#[...]` and `#![...]` which annotate the item they precede. The name of the test comes from the name of the function. Inside of test functions

there are a series of macros you can use for asserting facts about your system. We use `assert_eq` to ensure that the output of our `generate_sequence` function is what we expect it to be.

We can use `cargo` to run all of our tests:

```
$ cargo test
  Finished dev [unoptimized + debuginfo] target(s) in 0.43s
  Running target/debug/deps/numbers-f74640eac1a29f6d

running 1 test
test generate_sequence_should_work ... ok

test result: ok. 1 passed; 0 failed; 0 ignored; 0 measured; 0 filtered \
out
```

## Wrapping up

The Rust language is designed to allow fine grained control over performance and reliability while writing code at a high level of abstraction. The cost is learning the abstractions and dealing with the cognitive load of making choices as you design a program. These considerations are important for production quality code as you profile and optimize where you find bottlenecks. However, following the standards of the community will get you the majority of the way to high quality applications.

Our goal is to teach you Rust by building applications. This chapter is a bit of an anomaly as we took a long path to not get very far. We will try to avoid this as we progress by making pragmatic choices at the many points where it is possible to get bogged down in the small details about just the “right” way to do something. The standard library documentation is an excellent source to dig in to these details when you want to go deeper.

# Making A Web App With Actix

## Web Ecosystem

One area where Rust stands out is in the building of web servers.

Rust has its origins at Mozilla primarily as a tool for building a **new browser engine**. The existing engine being written in C++ combined with the syntactical similarities encourages the idea the Rust was meant to be a replacement for C++. There is obviously some truth to this, but in many ways this characterization sells Rust's potential short. While it is capable of being a systems programming language, there are a plethora of language features that make it suitable for innumerable programming tasks, including building web servers.

There are a few different layers to the web programming stack. Primarily we are concerned here with the application layer which is comparable to where Django, Rails, and Express live in Python, Ruby, and NodeJS, respectively.

The ecosystem around web application development in Rust is still quite nascent despite the Rust language hitting 1.0 in 2015. Much of the underlying infrastructure for building concurrent programs took until 2019 to reach a maturity sufficient to be included in the stable version of the standard library. However, most of the ecosystem has coalesced around similar ideas which take advantage of Rust's particular features.

Before jumping in to building a simple web server, let's briefly discuss a few of the libraries that make up the web landscape.

## Hyper


[Hyper](https://hyper.rs/)<sup>28</sup> is a low level HTTP library built on even lower level libraries for building network services. Currently most web frameworks use Hyper internally for handling the actual HTTP requests.

---

<sup>28</sup><https://hyper.rs/>

It can be used to build both HTTP clients and servers. However, there is a bit more boilerplate than you might want to write yourself when you want to focus on building an application. Therefore, we will use a library at a higher level of abstraction which still allows us to take advantage of what Hyper offers.

## Actix



Actix

The [Actix](https://actix.rs/)<sup>29</sup> project is actually a group of projects which define an actor system as well as a framework for building web applications. The web framework is aptly named `actix-web`. It has been built on top of futures and async primitives from the beginning. It also runs on the stable version of the compiler.

It recently hit the 1.0 milestone which should bring some much needed stability to the ecosystem. Additionally, it has been at the top of the [Tech Empower web framework benchmarks](https://www.techempower.com/benchmarks/#section=data-r16&hw=ph&test=plaintext)<sup>30</sup>. Even if those are artificial benchmarks, it still points to the performance potential possible.

Actix is the library that we are going to use in this chapter, but before we dive in, let's look at a few others from a high level.

---

<sup>29</sup><https://actix.rs/>

<sup>30</sup><https://www.techempower.com/benchmarks/#section=data-r16&hw=ph&test=plaintext>



# GET THE FULL BOOK



This is the end of the preview!

Get the full version at:

<https://newline.co/fullstack-rust>

Featuring:

- 7 lessons
- 240+ pages
- Complete code

 FULLSTACK.io

ANDREW WEISS

# GET IT NOW